Jan-Øivind Lima

Procedural generation: music

Generating music using method X and Y

Master's thesis in Cybernetics and Robotics
Supervisor: Sverre Hendseth
June 2026

Norwegian University of Science and Technology
Faculty of Information Technology and Electrical Engineering
Department of Engineering Cybernetics

NTNU

ABSTRACT

Write an abstract/summary of your thesis, and state your main findings here.

A summary should be included in both English and any second language, if this is applicable, regardless if the thesis is written in English or in your preferred language. These should be on separate pages, the English version first.

PREFACE

Write the preface of your thesis here.

You may include acknowledgements and thanks as part of your preface on this page, or you may add it as a new chapter after the preface.

This thesis is written as the final master project at Norwegian University of Science and Technology (NTNU) during the spring of 2026, supervised by Sverre Hendseth.

CONTENTS

Abstract **i**

Preface **ii**

Contents **iv**

List of Figures **iv**

List of Tables **v**

Abbreviations **vii**

1 Introduction **1**

2 Background **3**

2.1 Procedural generation 3

2.1.1 Wave function collapse 3

2.1.2 PG using constraint satisfaction 4

2.1.3 L-systems (Lindenmayer systems) 4

2.1.4 Binary space partitioning 4

2.1.5 Random noise and seeds 4

2.1.6 Random walks 4

2.2 Music theory 5

2.2.1 Notes 5

2.2.2 Scales 5

2.2.3 Keys 6

2.2.4 Melodies 6

2.2.5 Harmony 6

2.2.6 Time signatures 7

2.2.7 Musical form 7

2.3 Object oriented programming 9

2.3.1 Features of OOP 9

3 Software design **11**

3.1 Object oriented programming 11

4 Implementation **13**
4.1 The iteration process 13

5 Results **15**

6 Iterative development **16**
6.1 First iteration 16
6.1.1 Implementation 16
6.1.2 Results 17
6.1.3 Discussion 17
6.2 Second iteration 18

7 Discussion **19**
7.1 Future work 19
7.1.1 Non-block harmony 19

8 Conclusions **21**

References **23**

Appendices: **25**

A - Github repository **26**

B - Audio examples **27**

C - Code documentation **28**

LIST OF FIGURES

2.1.1 Illustration of how procedural tiles from input image is used to generate a new image. [2] 4
2.2.1 The c-major scale written in the treble clef. 5
2.2.2 The eb-major scale written in the treble clef. 6
2.2.3 Harmonic progression in tonal music. The arrows indicate the typical flow of harmonic functions. 7
2.2.4 A classical example of the musical sentence form. These are the opening bars of Bethovens piano sonata in F-minor. [3] 8

6.1.1 Melody showing how the octave is not restricted and ends up in unplayable range. 17
6.1.2 Melody showing how we can generate multiple voices with different rhythmic restrictions. 18
6.1.3 Example melody of the first iteration melody generator. 18

LIST OF TABLES

ABBREVIATIONS

List of all abbreviations in alphabetic order:

AI	Artificial Intelligence
LISP	List Processing
LLM	Large Language Models
MIT	Massachusetts Institute of Technology
NTNU	Norwegian University of Science and Technology
OOP	Object Oriented Programming
PG	Procedural Generation
WFC	Wave Function Collapse
QM	Quantum Mechanics

CHAPTER

ONE

INTRODUCTION

For hundreds of years composers have been creating music using external tools like dice to help them generate something new. As early as the 18th century the idea of using randomness to help the composing process has been used. The musical dice game, "Musikalisches Würfelspiel", was used to randomly generate music based on precomposed parts. [1] This is a very basic way of generating music, however by the random nature of the dice it will not be able to create anything coherent and meaningful like a human composer can. Additionally this process only used precomposed parts which were designed to work together. If we want to generate everything from scratch we need a way more sophisticated system. This is where procedural generation comes into the picute. If we base the meldoy on something random and then use the conventions and "rules" for melodic composition we might be able to generate music that sounds somewhat coherent.

Today the most common way to go about generating music is using Artificial Intelligence (AI) and Large Language Models (LLM). This enables anyone to create music, however these models are trained on large data sets and imitates what has come before. For this thesis the focus is on the craft of creating music, breaking it down to its very primitive form, only pitches and durations. This will enable the user to have more control over exactly what is being generated. However this requires us to create a system and models that capture the essence of how exactly we are to go about contructing the very basic musical ideas. It will not depend on large data sets and will create a musical ideas that can be processed into a fully fledged musical idea.

If we are able to solve this problem of creating a good structure and model for this kind of generation we will not only have a solution for musical ideas, but this kind of system can solve many kinds of problems.

The problem we want to solve is the problem of creating a coherent system that will be able to generate music that adheres to certain musical principles. In order to do this we need to consider many things, and system design is a big part of this. When buidling such a program we need to consider what parts it should consist of as well as how they are supposed to interact. Object Oriented Programming (OOP) is a solid way to approach this system design.

The solution will be a system that is modular, easy to modify and generate new music using newly defined constraints or inputs on structure, motifs etc...

CHAPTER

TWO

BACKGROUND

This chapter provides the background theory for the project. First we give an overview of procedural generation and some of the different methods used in procedural generation. Then we give an overview of some basic music theory that is necessary to fully grasp the structure of the system. Lastly we cover some about object oriented programming.

2.1 Procedural generation

Procedural Generation (PG) is algorithmic generation of content. This content can take many forms including landscape, levels, textures or even 3D models. This leads to smaller file sizes since the program can generate the content using algorithms rather than having everythign pre-made and saved.

This section will cover several methods of PG and discuss similarities, differences and most importantly limitations of the methods. Additionally we will take a look at and what we are trying to achieve with the procedural generation in this project and propose a new method that will handle the problem of music generation in a new way.

2.1.1 Wave function collapse

Wave Function Collapse (WFC) is a method used to procedurally generate using either given constraints, or it can also use some input that it analyses and generates the constraints procedurally which are then used to generate something in the same style as the input.

It works by initializing the entire space in a superposition which means that the state is yet to be determined. We call this an *unobserved* state, meaning we don't know what state it is in yet. These unobserved states are not complex numbers like they are in Quantum Mechanics (QM), but the process is similar which is why it gets its name inspired by the field of QM. Randomly assigning parts of the space we are able to collapse[1] some of the *wave function*[2]. In Figure 2.1.1 we

[1] collapsing a superposition means that we decide the actual state of the given position in the space

[2] the wave function is simply the constraints what decide what the state of the object should be dependent on external factors

Figure 2.1.1: Illustration of how procedural tiles from input image is used to generate a new image. [2]

can see an illustration of how the "tiles" of an input image re-appear in the final image generation.

2.1.2 PG using constraint satisfaction

Similar to WFC we have PG using constraint satisfaction. This method uses given constraints for influencing the generation. It usually works by randomly generating something in the given domain, then iterates over the generated data and enforces constraints in order to make it adhere to a model designed by the programmer or end user. After several iteration over the domain we end up with a domain where the constraints are satisfied for all the objects in the domain.

2.1.3 L-systems (Lindenmayer systems)

This is a way to generate things organically. The technique is often used in generating organic materials like treesh, bushes or other fractal-like systems.

2.1.4 Binary space partitioning

By partitioning the given space in smaller spaces recursively we are able to find a subdivision small enough to generate something that propagates to the higher space with certain characteristics.

2.1.5 Random noise and seeds

Perlin noinse is a very common way to generate structures like worlds etc... It is very good at creating smooth transitions from one state to another, like height maps.

2.1.6 Random walks

This works by randomly walking around a space and building a surface or hollow after that.

2.2 Music theory

This section contains some basic music theory relevant to understand how music is constructed. It is not neccessary to understand everything in depth, but farmiliarity with these concepts will be of great help to understand the project as a whole, not just the procedural generation part.

2.2.1 Notes

The very core of music is composed of notes. Notes are vibration in air which we interpret as pitches. In the western system we agree upon the standard of pitches is based around A4 which is $440Hz$. All other pitches are tuned in relation to this reference.

2.2.2 Scales

The musical scale is building block for creating not only melodies but also harmony, which we will cover a bit later. Firstly the musical scale is, for the most part, constructed using seven notes. In the western canon this is usually seven notes. You can see the basic c-major scale in Figure 2.2.1, it is comprised of the notes c d e f g a b c. The tones of the scale can be named according to their function within the scale. We often label the notes $\hat{1}, \hat{2}, \hat{3}, \hat{4}, \hat{5}, \hat{6}, \hat{7}$, and then starting back at $\hat{1}$. The use of the carets are called the Nashville Number System, this is used to differentiate the scale degrees from the chord tones which will be discussed in Section 2.2.5 When we consider chord structures we use the numbers but usually count above from 8 to signify harmonic functions, but when considering only scales and melodies we usually only use 1 through 7

The most important function is that of the scale degree 1, this is the *tonic*, which is the 'home' of the key. This is the most stable function there is, it feels resolved and does not have any tension urging it to move anywhere.

Next up we have the dominant function, there are two notes degrees which fall into this category, that is the 5th and 7th scale degree. These notes have tension. They want to move, especially the 7th scale degree, it has a great pull toward the tonic. You can hear this in the example under Figure 2.2.1. The penultimate tone is the 7th scale degree and it leads to the tonic.

To lead us to the tonic functions we have what are called the pre-dominant function. This simply means 'before dominant', and their function is to set up the tension to lead us to the dominant degrees which in turn leads us back to the tonic.

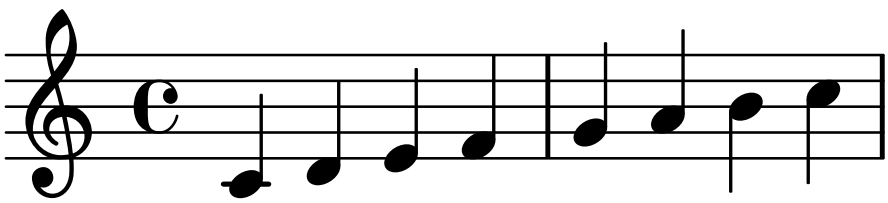

Figure 2.2.1: The c-major scale written in the treble clef.

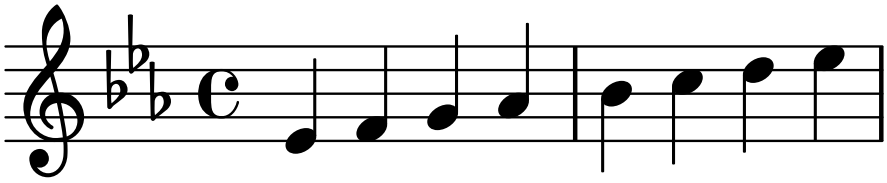

Figure 2.2.2: The eb-major scale written in the treble clef.

2.2.3 Keys

The concept of a musical key is simply shifting around where the tonic note is. For instance in Figure 2.2.1 you can see the major scale in c. And in Figure 2.2.2 you can see the major scale in eb. In todays tuning system there are no differences between the keys, meaning that the sound the same, just higher and lower. This was not the case earlier in history, where the temperament(how the tones were tunes in relation to each other) of the keys differed. Today it just sounds higher or lower depending on the key.

2.2.4 Melodies

Melodies is the horizontal component of music, meaning we look at it as a time-series. We have a given note at a given length followed by another note or rest of a given length etc... This means that you can look at a melody as a series of intervals as well. If we do this on known melodies we will observe that they are mainly composed of steps and skips. This means that the distance between one note and the next is one or two tones. Large leaps are usually leaping to important tones like the tonic or the root of the current chord.

2.2.5 Harmony

Harmony is the vertical component of the music, harmony is what happens when two or more notes happen at the same time. This is what chords are build around. When considering the harmonic progression in tonal music we can consider all the standard chords in the major key. They are labeled by numbers using the roman numerals. Upper case means the chord in major and lower case means minor. You can see the difference between major and minor chords in figure [make a figure for this], and listen to the difference at the link provided in the caption.

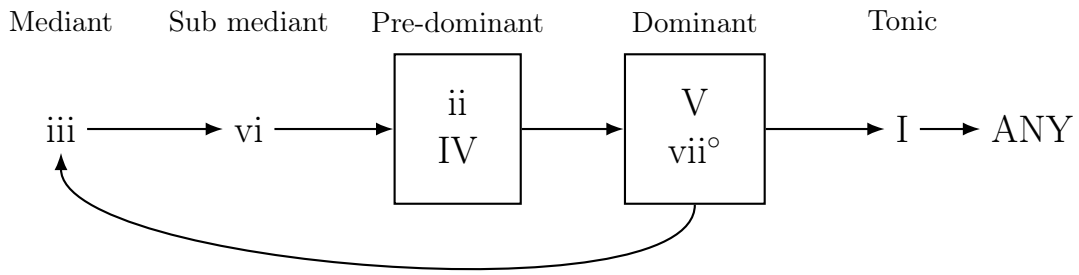

Figure 2.2.3: Harmonic progression in tonal music. The arrows indicate the typical flow of harmonic functions.

2.2.6 Time signatures

The time signature explains what is in a bar of music. A bar is the space between the vertical lines on the staff. A 4/4, sometimes written as "C", time means that there are four quarter notes per bar and consequently 3/4 time means that there are three quarter notes per bar.

2.2.7 Musical form

Now that we have an idea of the musical structures on the melodic and harmonic layer we will take a look at musical forms. This is a very important idea when it comes to constructing melodies. A melody is not an infinitely wandering idea of music. A good melody follows a form. Often this can be viewed as repetitions and variations of a given musical idea. Two very common forms of musical form are the *sentence form* and the *period form*.

In Figure 2.2.4 we can see an example of the musical sentence. It starts with a 2-bar opnening phrase, labeled as the *basic idea*, then we have a repetition of the same idea only adapted for a new chord. Then the following 2 bars are a fragmentation of the idea. Then for the last two bars we have a conclusion, this ties the idea together. The cadence at the end can be considered the end of a sentence. It concludes the idea before it moves on to something else.

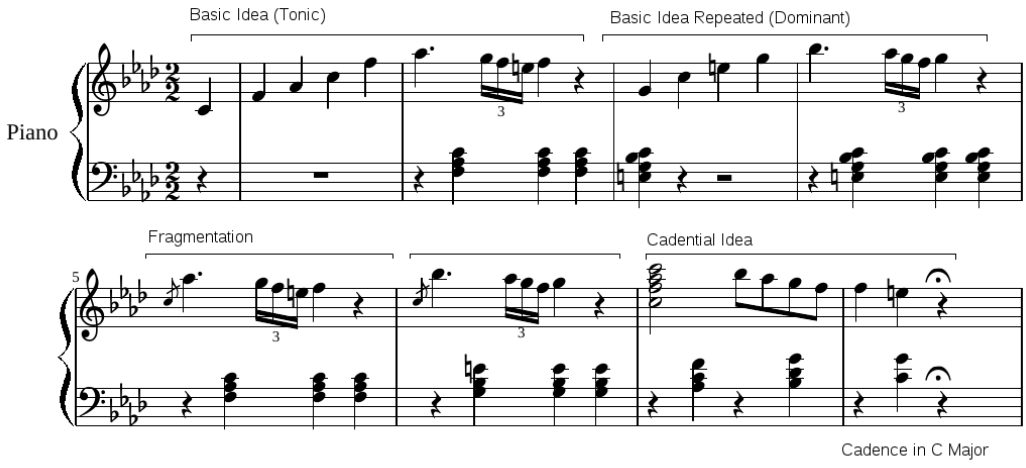

Figure 2.2.4: A classical example of the musical sentence form. These are the opening bars of Bethovens piano sonata in F-minor. [3]

2.3 Object oriented programming

OOP had its inception in the sixties with the *aritficial intelligence* group at Massachusetts Institute of Technology (MIT). Here the term "object" was used in the language family List Processing (LISP) to refer to LISP atoms[3] with with certain attributes or properties. [4]

2.3.1 Features of OOP

Some of the useful features of OOP are that we can achieve using OOP is ...

[3]a LISP atom is an indivisible entity in the language. LISP consist of two data types, atoms and lists. And a list can only contain atoms or other lists

CHAPTER

THREE

SOFTWARE DESIGN

Include the complete description of the methods used in your research here.

3.1 Object oriented programming

This project will use the widely known process of object oriented programming. In simple terms this simply means that all the parts of the code are structured as objects, meaning everything is represented as objects with given properties. This means that for instance a 'melody object' could consist of all the notes and durations it consists of, the key it is written in and any metadata explaining what the melody represents. Meaning if it is a verse or chorus etc...

In order to design an object oriented program the design process is very important. Creating good models for the objects are imperative to make it work. If the level of abstraction is too high, the sematic gap between the real-world concepts they represent and the acutal implementation in the code will be too large and the model will be un-intuitive and hard to both use and also understand.

CHAPTER

FOUR

IMPLEMENTATION

The implementation of this program started by just writing something that generated something. This was done to not get stuck in the planning phase. You could consider this a contradiction to the part in software design which states that good planning is imperative. However, trying out some code and figuring out what works, what doesn't and what needs to be implemented in the models is very helpful. Therefore the first iteration was a test to get some systems working and finding out what it neccesary to make it work and what is not imporatant.

4.1 The iteration process

After the first iteration lots of issues were observed and were put into the spec for the second iteration which improved upon the model in various ways. Some imporatnt features were restricting the range so that an acutal instrument can play it, be that an instrument or the human voice.

CHAPTER

FIVE

RESULTS

CHAPTER

SIX

ITERATIVE DEVELOPMENT

In this project I have utilized a from of development called iterative design. This means that I do the project by creating a design, taking notes on what needs to be improved, then starting again with the new knowledge gained in the previos iteration.

This sounds like more work, but in reality it leads to an effective workflow. It allows for fast prototyping since we can start programming something straight away, then when the project halts due to the system that has been created is incomptabile with the ideas for improving it we take notes on what worked, what is missing, and use this to create a specification for the next iteration.

By doing this we maintain active development which doesn't stagnate and we can always create a new iteration when development slows down. After 3-4 iterations we should be able to land on a system that works in a way that we are happy with.

6.1 First iteration

The goal of the first iteration was to get a "proof of concept" working. I wanted to test out how I could generate notes and visualize them as well as listen to audio of the notes generated. For this I used *python* to generate *liypond* files. From these files I was able to generate PDFs of sheet music as well as *wave-files* to listen to how it sounds without having to play/sing it myself.

6.1.1 Implementation

This first iteration was a very primitive handling of generating music. It had very few constraints on the melody and contained little to no metadata on the music which made it unsuitable for actual music.

The code generated the notes within a given scale and restricted the distance (called interval in music theory) to only 2nds and 3rds. However this meant that the melody meandered without any other restrictions. If the generated music was long it would have a tendency to go beyond the range of an instrument or the human voice.

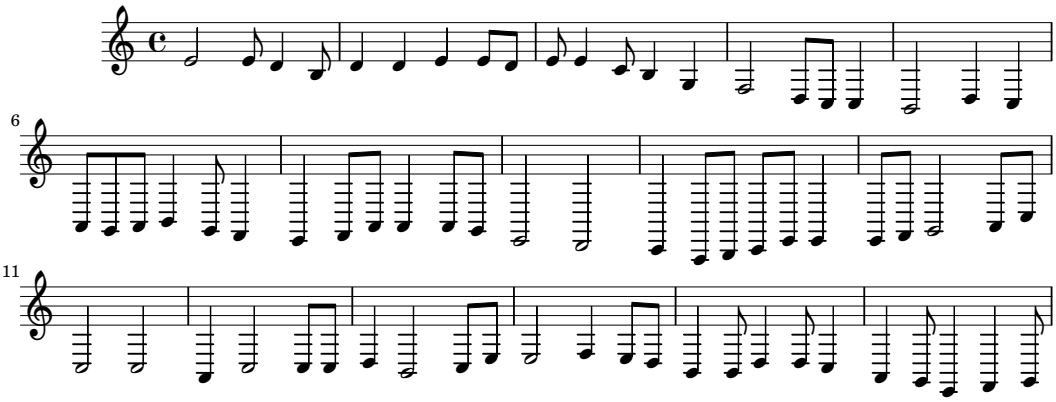

Figure 6.1.1: Melody showing how the octave is not restricted and ends up in unplayable range.

6.1.2 Results

The melody generation works fine for smaller melodies, however there is no self-similarity and no rests. This means that the melodies we generate are not particularly musical or interesting, but it is a great starting point to getting to a fully fledged melody generator. In Figure 6.1.3 we can see how a short melody using the first iteration looks.

In Figure 6.1.1 you can see what happens when the generated melody is longer, it will tend to go beyond what an instrument will be able to play. You can see this clearly in *bar 6 through 16* (the entire 2nd and 3rd system), where there are lots of ledger lines. Some of these notes are not even available on the grand piano.

However a good feature that was included was the ability to generate more voices at the same time. This feature was quite primitive since the voices were generated independently and did not adhere to any harmony, however it can be a good skeleton to use later if we get to the point that we want to actually harmonize with the melody we generate. In Figure 6.1.2 you can see how we generated four voices all with their own rhythmic restriction. The bottom voice only has whole notes, the second lowest one has only half-notes, the second highest has only quarter and the top voice has a combination of several types.

6.1.3 Discussion

As we can see the first iteration shows a promising proof of concept, however it is quite limited and needs lots of improvement before we have a good system for melody generation.

For the second iteration we need to be able to enforce more restictions on the melody so that it follows more convetions of good melody writing. Additionally we need to have some self-similarity so that the music follows some kind of musical form in order for it to feel coherent and part of the same idea.

The second iteration will focus on creating a generation algorithm that tat is able to generate a melody with a structure. Meaning when generating a melody we will generate a phrase, that phrase should contain an idea, a repetition, and the end of the melody should end resolved with a long note on the tonic of the

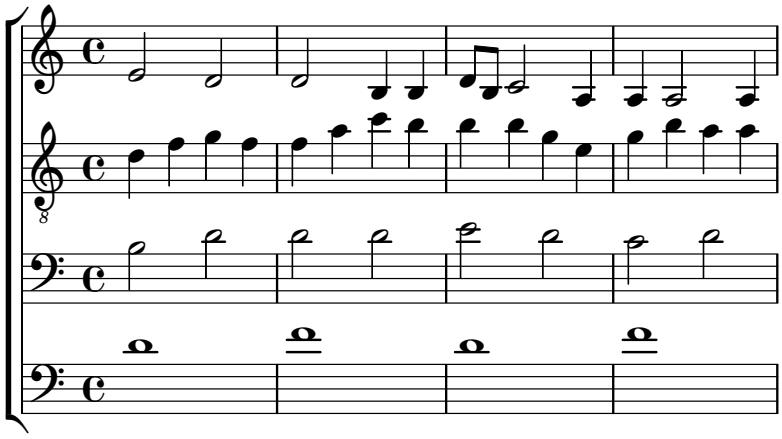

Figure 6.1.2: Melody showing how we can generate multiple voices with different rhythmic restrictions.

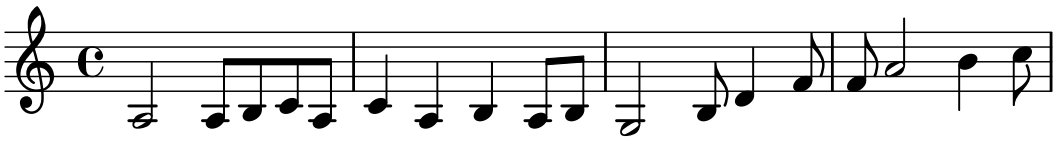

Figure 6.1.3: Example melody of the first iteration melody generator.

key. Additionally we need to restirct the range of the melody so that it will be playable/singable.

6.2 Second iteration

CHAPTER

SEVEN

DISCUSSION

Discuss your results here.

The generated music is good because this:

The generated music is bad because this:

etc etc...

7.1 Future work

The current implementation is not complete and there are several features which would benefit the system.

7.1.1 Non-block harmony

Not sure how far the project will go, but probably no implementatino of Non-block harmony happened.

CHAPTER

EIGHT

CONCLUSIONS

Give a concise summary of your research and finding here, and include a short summary of any future work as well.
Looking at this thing here, [5] [6] [7] it is a citation.

REFERENCES

[1] *Musikalisches Würfelspiel*. In: *Wikipedia*. Sept. 27, 2025. URL: https://en.wikipedia.org/w/index.php?title=Musikalisches_W%C3%BCrfelspiel&oldid=1313693497 (visited on 03/24/2026).

[2] Maxim Gumin. *Wave Function Collapse Algorithm*. Version 1.0. Sept. 2016. URL: https://github.com/mxgmn/WaveFunctionCollapse (visited on 03/24/2026).

[3] *Sentence (Music)*. In: *Wikipedia*. June 21, 2023. URL: https://en.wikipedia.org/w/index.php?title=Sentence_(music)&oldid=1161220566 (visited on 03/26/2026).

[4] *Object-Oriented Programming*. In: *Wikipedia*. Mar. 23, 2026. URL: https://en.wikipedia.org/w/index.php?title=Object-oriented_programming&oldid=1344935717 (visited on 03/27/2026).

[5] Roland van der Linden, Ricardo Lopes, and Rafael Bidarra. "Procedural Generation of Dungeons". In: *IEEE Transactions on Computational Intelligence and AI in Games* 6.1 (Mar. 2014), pp. 78–89. ISSN: 1943-0698. DOI: 10.1109/TCIAIG.2013.2290371. URL: https://ieeexplore.ieee.org/abstract/document/6661386 (visited on 02/26/2026).

[6] E. Galin et al. "Procedural Generation of Roads". In: *Computer Graphics Forum* 29.2 (2010), pp. 429–438. ISSN: 1467-8659. DOI: 10.1111/j.1467-8659.2009.01612.x. URL: https://onlinelibrary.wiley.com/doi/abs/10.1111/j.1467-8659.2009.01612.x (visited on 02/26/2026).

[7] Johan Gangsås Hole. "Automatic Species Counterpoint". In: *Master Thesis at NTNU* (2026).

APPENDICES

A - GITHUB REPOSITORY

All code and latex-files used in this document are included in the Github repository linked below. Further explanations are given in the readme-file.

Github repository link

- `https://github.com/lima98/TTK4900`

B - AUDIO EXAMPLES

Audio-files for the examples in the thesis can be played from my NTNU-webpage.

Audio example archive

- `https://www.janoivil.folk.ntnu.no`

C - CODE DOCUMENTATION

All documentation for the code is included in the Github repository linked in 8 it is also provided as a web-page hosted at my NTNU-webpage.

Code documentation

- `https://janoivil.folk.ntnu.no/docs/_build/html/index.html`